



MACHINE LEARNING

Soheil Shirvani 810195416



Question 1:

Subject: _____ Date: _____

سوال ۱:

ما در این شکل می بینیم که فصل ها توسط خط $y = -x$ قرار دارد.

و نقاط $(0, 1)$ روی $y = -x + 1$ و نقاط $(1, 0)$ و $(-1, 0)$ روی $y = -x - 1$ است.

با توجه به رابطه فصل مشترک داریم: $h(\vec{x}) = w_1 x + w_2 y + b \geq 0$

$\Rightarrow y < -x \Rightarrow -x - y \geq 0 \Rightarrow w_1 = -1, w_2 = -1, b = 0$

می دانیم در SVM ها scale نداریم پس $w_1 = -1, w_2 = -1, b = 0$ نیز می تواند درست باشد.

هر عدد دلخواه است.

می دانیم بهترین فاصله با استفاده از این وزن ها برابر

$$\frac{2}{\|w\|} = \frac{2}{\sqrt{(-1)^2 + (-1)^2}} = \frac{2}{\sqrt{2}}$$

و از روی شکل می بینیم بهترین فاصله ی بین نقاط برابر ۲ است پس داریم:

$$\frac{2}{\sqrt{2}} = 2 \Rightarrow C = \frac{\sqrt{2}}{2}$$

$\Rightarrow w_1 = -\frac{\sqrt{2}}{2}, w_2 = -\frac{\sqrt{2}}{2}, b = 0$

و $SV = \{(0, 1), (1, 0), (-1, 0)\}$ و معادله جدا سازی: $-\frac{\sqrt{2}}{2}x - \frac{\sqrt{2}}{2}y \geq 0$

Question 2:

1)

Solving the primal problem, we obtain the optimal w , but **know nothing about the** α_i . In order to classify a query point x we need to explicitly compute the scalar product $w^T x$, which may be **expensive** if d is large.

Solving the dual problem, we obtain the α_i (where $\alpha_i=0$ for all but a few points - the support vectors). In order to classify a query point x , we calculate

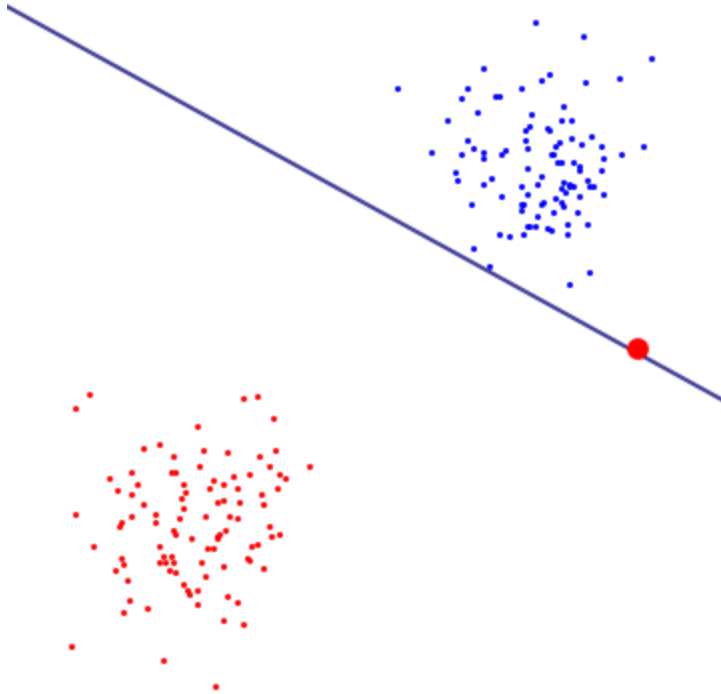
$$w^T x + w_0 = \left(\sum_{i=1}^n \alpha_i y_i x_i \right)^T x + w_0 = \sum_{i=1}^n \alpha_i y_i \langle x_i, x \rangle + w_0$$

This term is very **efficiently calculated** if there are only few support vectors. Further, since we now have a scalar product only involving *data* vectors, we may **apply the kernel trick**.

2)

I would expect soft-margin SVM to be better even when training dataset is linearly separable. The reason is that in a hard-margin SVM, a single outlier can determine the boundary, which makes the classifier overly sensitive to noise in the data.

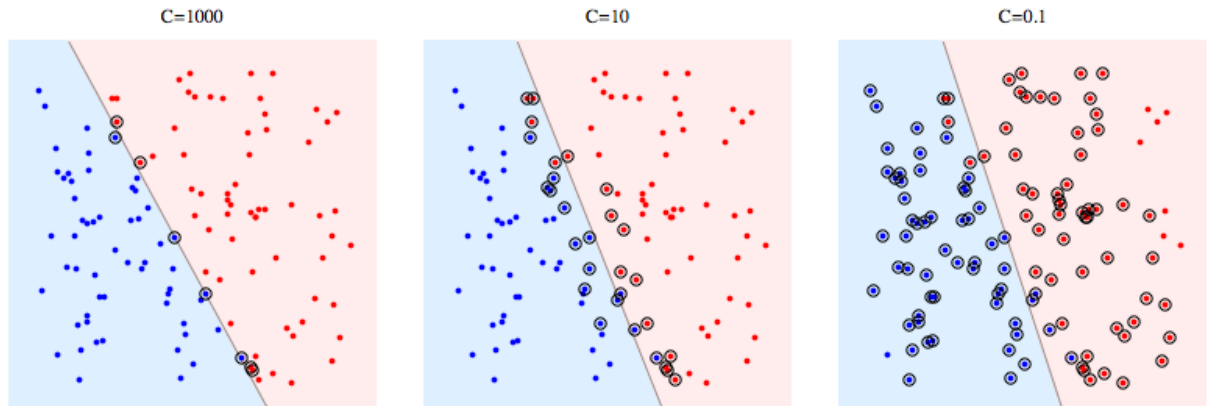
In the diagram below, a single red outlier essentially determines the boundary, which is the hallmark of overfitting



To get a sense of what soft-margin SVM is doing, it's better to look at it in the dual formulation, where you can see that it has the same margin-maximizing objective (margin could be negative) as the hard-margin SVM, but with an additional constraint that each lagrange multiplier associated with support vector is bounded by C . Essentially this bounds the influence of any single point on the decision boundary, for derivation, see Proposition 6.12 in Cristianini/Shaw-Taylor's "An Introduction to Support Vector Machines and Other Kernel-based Learning Methods".

The result is that soft-margin SVM could choose decision boundary that has non-zero training error even if dataset is linearly separable, and is less likely to overfit.

Here's an example using libSVM on a synthetic problem. Circled points show support vectors. You can see that decreasing C causes classifier to sacrifice linear separability in order to gain stability, in a sense that influence of any single datapoint is now bounded by C .



Meaning of support vectors:

For hard margin SVM, support vectors are the points which are "on the margin". In the picture above, $C=1000$ is pretty close to hard-margin SVM, and you can see the circled points are the ones that will touch the margin (margin is almost 0 in that picture, so it's essentially the same as the separating hyperplane)

For soft-margin SVM, it's easier to explain them in terms of dual variables. Your support vector predictor in terms of dual variables is the following function.

$$\begin{aligned}
 f(\mathbf{x}, \boldsymbol{\alpha}^*, b^*) &= \sum_{i=1}^{\ell} y_i \alpha_i^* \langle \mathbf{x}_i \cdot \mathbf{x} \rangle + b^* \\
 &= \sum_{i \in \text{sv}} y_i \alpha_i^* \langle \mathbf{x}_i \cdot \mathbf{x} \rangle + b^*.
 \end{aligned}$$

Here, alphas and b are parameters that are found during training procedure, \mathbf{x}_i 's, y_i 's are your training set and \mathbf{x} is the new datapoint. Support vectors are datapoints from training set which are included in the predictor, ie, the ones with non-zero alpha parameter.

3)

The C parameter tells the SVM optimization how much you want to avoid misclassifying each training example. For large values of C , the optimization will choose a smaller-margin hyperplane if that hyperplane does a better job of getting all the training points classified correctly. Conversely, a very small value of C will cause the optimizer to look for a larger-margin separating hyperplane, even if that hyperplane misclassifies more points. For very tiny values of C , you should get misclassified examples, often even if your training data is linearly separable.

Question 3:

Kernel Function is a method used to take data as input and transform into the required form of processing data. "Kernel" is used due to set of mathematical functions used in Support Vector Machine provides the window to manipulate the data. So, Kernel Function generally transforms the training set of data so that a non-linear decision surface is able to transformed to a linear equation in a higher number of dimension spaces. Basically, it returns the inner product between two points in a standard feature dimension.

Falat Pejvak Co.
 Inspection & Certification

Subject: _____ Date: _____

$$|K(x, y)|^2 \leq K(x, x) K(y, y) \Rightarrow \langle x, z \rangle^2 \leq \|x\|^2 \|z\|^2$$

$$\Rightarrow 0 \leq \|(\|z\| + \epsilon)x \pm z(\|x\| + \epsilon)\|^2$$

$$= \langle (\|z\| + \epsilon)x \pm z(\|x\| + \epsilon), (\|z\| + \epsilon)x \pm z(\|x\| + \epsilon) \rangle$$

$$= (\|z\| + \epsilon)^2 \|x\|^2 + \|z\|^2 (\|x\| + \epsilon)^2 \pm 2 \langle (\|z\| + \epsilon)x, z(\|x\| + \epsilon) \rangle$$

$$\leq 2(\|z\| + \epsilon)^2 (\|x\| + \epsilon)^2 \pm 2(\|z\| + \epsilon)(\|x\| + \epsilon) \langle x, z \rangle$$

$$\Rightarrow \pm \langle x, z \rangle \leq (\|x\| + \epsilon)(\|z\| + \epsilon) \xrightarrow{\epsilon \rightarrow 0} \langle x, z \rangle^2 \leq \|x\|^2 \|z\|^2$$

$$Q(x, x) = \sum_{i=1}^Q \sum_{j=1}^Q a_i a_j K(x_i, x_j) \rightarrow E[\varphi(x)] = E[\dots]$$

$$= \sum_{i=1}^Q \sum_{j=1}^Q E[a_i a_j] K(x_i, x_j) = \frac{1}{Q^2} \sum_{i=1}^Q \sum_{j=1}^Q K(x_i, x_j) = E(\|x\|^2) = \|x\|^2$$

$$\frac{1}{Q^2} = \text{mean of each coefficient twice}$$

$$\Rightarrow \|x\| = \sqrt{\frac{1}{Q^2} \sum_{i=1}^Q \sum_{j=1}^Q K(x_i, x_j)} = \frac{1}{Q} \sqrt{\sum_{i=1}^Q \sum_{j=1}^Q K(x_i, x_j)}$$

Question 4:

(4) سوال

Subject: Date:

Falat Pejvak Co.
Inspection & Certification

① $k(x, y) = f(x) k_1(x, y) f(y)$, (چون $\phi: \mathcal{X} \mapsto f(x) \in \mathbb{R} (1d)$)
 $\Rightarrow f(x) f(y) = k(x, y) \Rightarrow f(x) f(y) \rightarrow \text{Valid kernel}$, 4 ویژگی: k_1, k_2, \dots
 $\Rightarrow f(x) \cdot k_1(x, y) \cdot f(y) \rightarrow \text{Valid kernel}$

② $k(x, y) = \exp(k_1(x, y)) \rightarrow 1e^{x^2} = \sum_{n=0}^{\infty} \frac{x^n}{n!} \Rightarrow e^{k_1} = \sum_{n=0}^{\infty} \frac{k_1^n}{n!}$
 $e^{k_1} = 1 + \frac{k_1^1}{1} + \dots$, ویژگی: $k_1 \times k_2 \rightarrow \text{Valid kernel}$
 $k_1 + k_2 \rightarrow \text{"}$, $a \cdot k \rightarrow \text{"}$ $\Rightarrow e^{k_1} \rightarrow \text{Valid kernel}$

③ $k(x, y) = k_1(x, y) + k_2(x, y)$
 چون $\phi(x) = [\phi_1(x), \phi_2(x)]$, $k(x, y) = \langle \phi(x), \phi(y) \rangle$
 $= \langle [\phi_1(x), \phi_2(x)], [\phi_1(y), \phi_2(y)] \rangle = \langle \phi_1(x), \phi_1(y) \rangle + \langle \phi_2(x), \phi_2(y) \rangle$
 $= k_1(x, y) + k_2(x, y)$

④ $k(x, y) = k_1(x, y) k_2(x, y)$
 چون $\phi(x)_i = \phi_1(x)_i \cdot \phi_2(x)_i$, $k(x, y) = \langle \phi(x), \phi(y) \rangle$
 $= \sum_{i=1}^m \sum_{j=1}^n \phi(x)_i \phi(y)_j = \sum_{i=1}^m \phi_1(x)_i \phi_1(y)_i \sum_{j=1}^n \phi_2(x)_j \phi_2(y)_j = k_1(x, y) k_2(x, y)$

⑤ $k(x, y) = x^T A y$, $A = \text{semidefinite} \rightarrow A = V \Lambda V^T$, $\Lambda = \text{مقادیر}$
 $k(x, y) = x^T A y = x^T V \Lambda V^T y$, چون $B = \sqrt{\Lambda} V^T$
 $\Rightarrow x^T V \Lambda V^T y = x^T B^T B y = \langle Bx, By \rangle \rightarrow \text{طبیعی}$
 $\Rightarrow \text{Valid kernel}$

Question 7: Coding:

In this part we are going to change hyper parameters of a 2-layer perceptron and compare the parameters. It has a input layer, hidden layer, and an output layer

Our parameters are:

```
1 hidden_layer_size = "[100, 50]"
2 activation_function = "[relu, sigmoid]"
3 optimizer = "[adam, sgd]"
4 max_iteration = "[20, 30]"
5 early_stopping = "[True, False]"
```

First we are going to discuss about each individual parameter here:

1) Hidden Layer Size:

This parameter defines number of neurons in the hidden layer. We can increase or decrease the size of the hidden layer using number of neurons in our code.

2) Activation Function:

- Relu: The rectified linear activation function or Relu for short is a piecewise linear function that will output the input directly if it is positive, otherwise, it will output zero. ... The rectified linear activation function overcomes the vanishing gradient problem, allowing models to learn faster and perform better. $F(x) = x ; x \geq 0$
- Sigmoid: In sigmoid activation, the input to the function is transformed into a value between 0.0 and 1.0. Inputs that are much larger than 1.0 are transformed to the value 1.0, similarly, values much smaller than 0.0 are snapped to 0.0. Shape of the function for all possible inputs is an S-shape from 0 up through 0.5 to 1.0. $F(x) = \frac{1}{1 + e^{-x}}$

3) Optimizer:

- Adam: Adam is a replacement optimization algorithm for stochastic gradient descent for training deep learning models. Adam combines the best properties of the AdaGrad and RMSProp algorithms to provide an optimization algorithm that can handle sparse gradients on noisy problems. Adam (short for Adaptive Moment Estimation) is an update to the RMSProp optimizer. In this optimization algorithm, running averages of both the gradients and the second moments of the gradients are used.
- SGD: Stochastic gradient descent (often abbreviated SGD) is an iterative method for optimizing an objective function with suitable smoothness properties. classical stochastic gradient descent is generally sensitive to the step size η . Fast convergence requires large step sizes but this may induce numerical instability. The problem can be largely solved[13] by considering implicit updates whereby the stochastic gradient is evaluated at the next iterate rather than the current one.

4) Max Iteration:

Max Iteration is the number of epochs which our model will run before ending. More epochs mean more learning and the overfitting will be increased.

5) Early Stopping:

stopping the training of neural network early before it has overfit the training dataset can reduce overfitting and improve the generalization of deep neural networks. The challenge of training a neural network long enough to learn the mapping, but not so long that it overfits the training data. Model performance on a holdout validation dataset can be monitored during training and training stopped when generalization error starts to increase. The use of early stopping requires the selection of a performance measure to monitor, a trigger for stopping training, and a selection of the model weights to use.

Our Default model is Default Network: Optimizer=Adam, NumberOfNeuron=100, Activation=Relu, EarlyStopping=False, Epoch=20, And for comparing we just change that specific parameter and save all the rest. For the next part, we then write back our default network and change the parameters with respect to the question again.

Now We can see The Result of each change here:

1. Changing number of neurons in hidden layer:

1) 50:

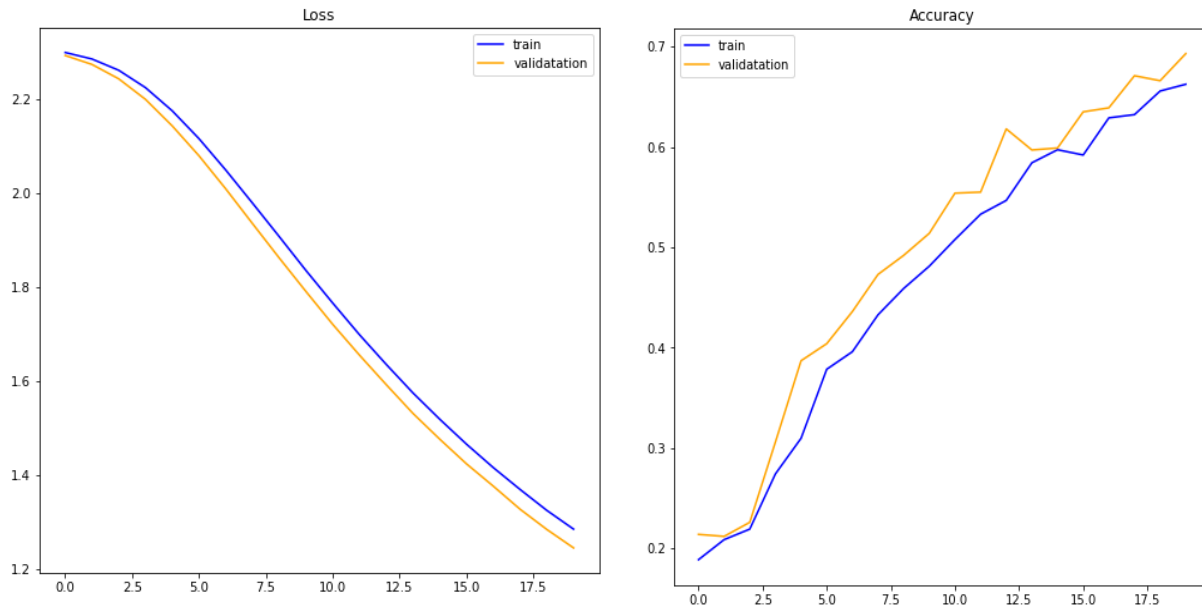
Epoch 19/20

125/125 [=====] - 0s 2ms/step - loss: 1.3242 - accuracy: 0.6557 - val_loss: 1.2837 - val_accuracy: 0.660

Epoch 20/20

125/125 [=====] - 0s 2ms/step - loss: 1.2843 - accuracy: 0.6625 - val_loss: 1.2444 - val_accuracy: 0.6930

79/79 [=====] - 0s 1ms/step - loss: 1.3651 - accuracy: 0.6180



Test Accuracy is : 61.800

Train Accuracy is : 66.250

Train Confusion Matrix

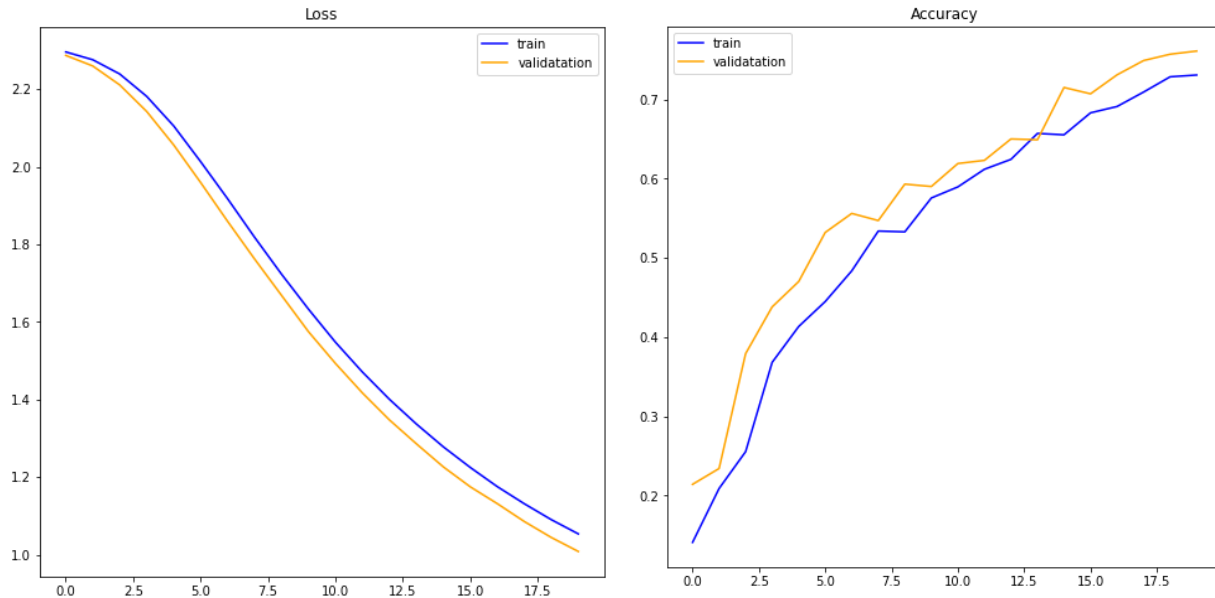
```
[[474  0  2  22  2  1  19  0  2  0]
 [  0 552  3  7  0  0  0  2  0  0]
 [ 11 43 374 31  5  0  32  3  3  1]
 [ 14 41 14 433  1  0  10 18  4  4]
 [  7 22  2  0 250  0  13 13  2 136]
 [ 32 48  4 195 15 31 19 33 46  6]
 [  9 23 25  4  3  0 448  0  1  0]
 [  4 51  0  2  7  0  1 437  0 18]
 [ 27 92 27 153  5  0 17  4 136 21]
 [ 13 22  1  5 56  0  2 152  1 231]]
```

Test Confusion Matrix

```
[[199  0  0  5  0  1 13  0  1  0]
 [  0 281  1  3  0  0  2  0  0  0]
 [  9 39 167 30  2  0 22  6  0  1]
 [  1 20 10 201  1  0  2 14  1  4]
 [  6 12  3  0 121  0 12 16  0 105]
 [ 16  5  6 111  4 10 14 31 15  9]
 [  7 13 12  5  4  0 183  0  0  1]
 [  1 38  2  1  4  0  0 199  3  9]
 [ 11 33 12 91  6  0  5  9 56 19]
 [  3 16  1  4 31  0  1 58  2 128]]
```

2) 100 Neurons:

Epoch 19/20
 125/125 [=====] - 0s 1ms/step - loss: 1.0902 - accuracy: 0.7285 - val_loss: 1.0442 - val_accuracy: 0.7570
 Epoch 20/20
 125/125 [=====] - 0s 2ms/step - loss: 1.0535 - accuracy: 0.7308 - val_loss: 1.0082 - val_accuracy: 0.7610
 79/79 [=====] - 0s 1ms/step - loss: 1.1392 - accuracy: 0.7092



Test Accuracy is : 70.920

Train Accuracy is : 73.075

Train Confusion Matrix

```
[[479  0  7  7  2  8 11  0  8  0]
 [  0 548  5  2  0  0  0  2  7  0]
 [  8 35 413 12  7  0 19  3  6  0]
 [  6 23 29 434  2  5  3 17 17  3]
 [  1 11  1  0 359  0 15 10  0 48]
 [ 28 46 17 116 24 105 12  8 55 18]
 [  7 21 14  0  3  2 464  0  2  0]
 [  3 33  2  0  8  0  0 452  1 21]
 [ 21 63 41 68 11  1 14  5 240 18]
 [ 12  9  2  5 86  0  1 151  2 215]]
```

Test Confusion Matrix

```
[[201  0  1  1  0  2 10  0  4  0]
 [  0 279  1  1  0  0  2  0  4  0]
 [  9 23 207  9  2  0 16  7  3  0]
 [  0  7 15 209  0  1  1 14  5  2]
 [  1  8  1  0 194  0 15  8  0 48]
 [ 15  4 11 81 11 51  7 15 14 12]
 [  6 11  9  1  4  2 192  0  0  0]
 [  1 26  3  0  3  0  0 214  3  7]
 [ 10 19 28 39 10  0  3 11 109 13]
 [  3  7  1  3 56  1  1 51  4 117]]
```

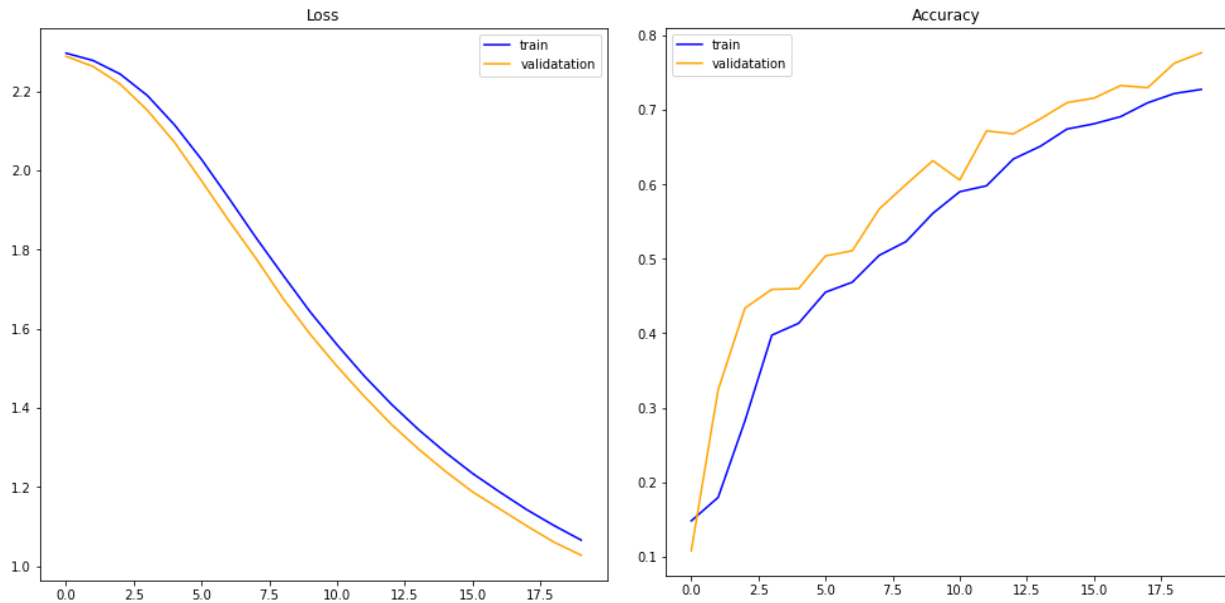
Here we can see by increasing the number of neurons to 100, a better accuracy was achieved.

There is no particular way to determine what number of neurons is better for the model, we should find it by experiments. In this question since our input shape had 196 feature it was kind of guessing that 50 is small. It was proved by experiment as 100 neurons achieved a better accuracy.

2. Activation Functions:

1) Relu:

```
Epoch 19/20
125/125 [=====] - 0s 2ms/step - loss: 1.1031 - accuracy: 0.7222 - val_loss: 1.0611 - val_accuracy: 0.7
630
Epoch 20/20
125/125 [=====] - 0s 2ms/step - loss: 1.0664 - accuracy: 0.7278 - val_loss: 1.0282 - val_accuracy: 0.7
770
79/79 [=====] - 0s 1ms/step - loss: 1.1662 - accuracy: 0.7060
```



Test Accuracy is : 70.600

Train Accuracy is : 72.775

Train Confusion Matrix

```
[[485  0  2  5  4  8 10  0  8  0]
 [ 0 548  4  4  0  4  0  1  2  1]
 [ 12 31 388 14  9  0 37  3  9  0]
 [ 11 24 22 427  5  8  4 14 15  9]
 [ 2  6  0  0 339  0 15  2  0 81]
 [ 22 33  3 129 24 136 15 12 29 26]
 [ 5 19  6  0  3  2 475  0  3  0]
 [ 4 24  1  0  8  0  1 436  5 41]
 [ 27 56 20 82  6  0 19  5 241 26]
 [ 12  9  0  5 62  0  3 56  1 335]]
```

Test Confusion Matrix

```
[[202  0  0  0  2  4  9  0  2  0]
 [ 0 279  0  3  0  0  3  0  2  0]
 [ 10 19 192 10  2  1 27  7  7  1]
 [ 0  8 10 205  0  4  3 12  6  6]
 [ 1  6  0  0 189  0 14  1  0 64]
 [ 12  4  6 89  8 45 10 21  8 18]
 [ 7  9  2  1  6  1 199  0  0  0]
 [ 1 19  5  0  4  0  0 188  4 36]
 [ 9 20 14 51  9  0  4  8 103 24]
 [ 3  7  1  4 47  0  1 15  3 163]]
```

2) Sigmoid:

Epoch 19/20

125/125 [=====] - 0s 2ms/step - loss: 2.2094 - accuracy: 0.3185 - val_loss: 2.1976 - val_accuracy: 0.5

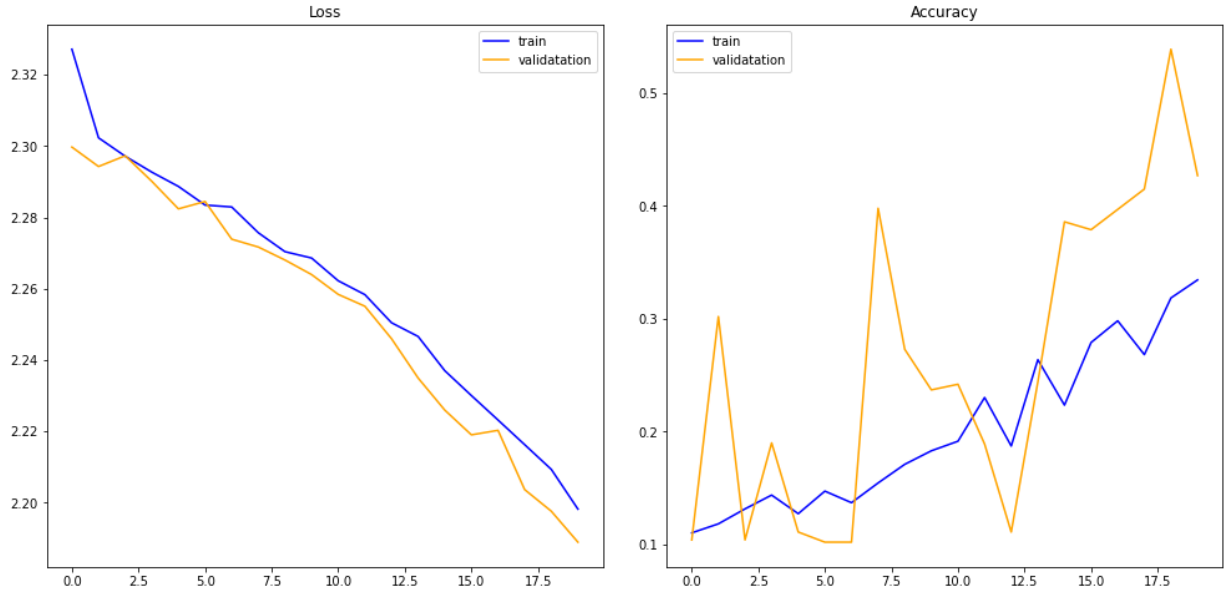
390

Epoch 20/20

125/125 [=====] - 0s 1ms/step - loss: 2.1982 - accuracy: 0.3345 - val_loss: 2.1889 - val_accuracy: 0.4

270

79/79 [=====] - 0s 1ms/step - loss: 2.2143 - accuracy: 0.3808: 0s - loss: 2.2136 - accuracy: 0.38



Test Accuracy is : 38.080

Train Accuracy is : 33.450

Train Confusion Matrix

```

[[473  2  0 46  0  0  1  0  0  0]
 [  0 554  0 10  0  0  0  0  0  0]
 [ 83 116 44 219  0  0 41  0  0  0]
 [ 10  36  0 489  0  0  3  0  0  1]
 [ 68 135  0 104  0  0 39  0  0 99]
 [ 34 133  0 258  0  0  2  0  0  2]
 [ 38  70  0  22  0  0 383  0  0  0]
 [ 45 269  0  89  0  0  1  0  0 116]
 [ 43 143  0 287  0  0  6  0  0  3]
 [ 51 158  0 101  0  0  0  0  0 173]]

```

Test Confusion Matrix

```

[[202  0  0 16  0  0  1  0  0  0]
 [  0 279  0  7  0  0  1  0  0  0]
 [ 20  82 13 144  0  0 16  0  0  1]
 [  1  15  0 237  0  0  0  0  0  1]
 [ 45  78  0  87  0  0 16  0  0 49]
 [ 16  32  0 171  0  0  1  0  0  1]
 [ 28  37  0  16  0  0 144  0  0  0]
 [ 11 163  0  59  0  0  0  0  0 24]
 [ 13  47  0 176  0  0  1  0  0  5]
 [ 12  69  0  85  0  0  1  0  0 77]]

```

Here we test 2 different activation functions Relu and Sigmoid.

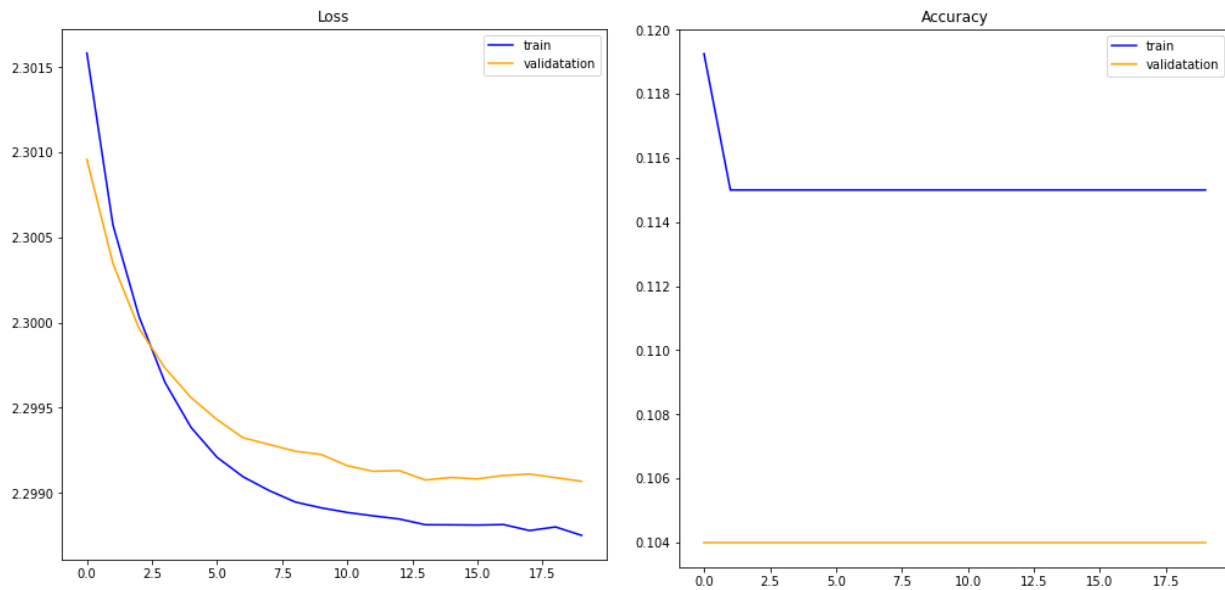
As we can see, model using sigmoid could not fit on the dataset and its loss is noisy. Its accuracy is low too. Sigmoid have a lot of problem including climbing, gradient vanishing and etc. It is not a good activation function; it was used in the passed before the appearance of new activations like Relu.

Relu achieved a great accuracy and it is good activation for general purposes and will often result in good accuracy.

3. Optimizer:

1) SGD

```
Epoch 19/20
125/125 [=====] - 0s 2ms/step - loss: 2.2988 - accuracy: 0.1150 - val_loss: 2.2991 - val_accuracy: 0.1040
Epoch 20/20
125/125 [=====] - 0s 2ms/step - loss: 2.2988 - accuracy: 0.1150 - val_loss: 2.2991 - val_accuracy: 0.1040
79/79 [=====] - 0s 1ms/step - loss: 2.3020 - accuracy: 0.1148
```



Test Accuracy is : 11.480

Train Accuracy is : 11.500

Train Confusion Matrix

```
[[ 0 522  0  0  0  0  0  0  0  0]
 [ 0 564  0  0  0  0  0  0  0  0]
 [ 0 503  0  0  0  0  0  0  0  0]
 [ 0 539  0  0  0  0  0  0  0  0]
 [ 0 445  0  0  0  0  0  0  0  0]
 [ 0 429  0  0  0  0  0  0  0  0]
 [ 0 513  0  0  0  0  0  0  0  0]
 [ 0 520  0  0  0  0  0  0  0  0]
 [ 0 482  0  0  0  0  0  0  0  0]
 [ 0 483  0  0  0  0  0  0  0  0]]
```

Test Confusion Matrix

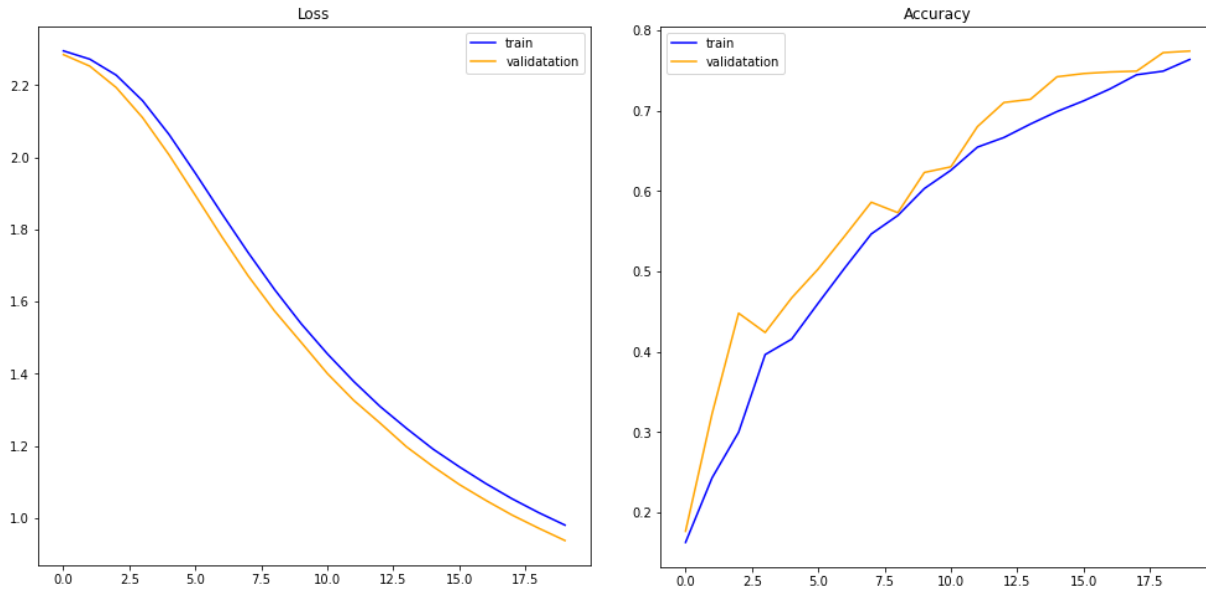
```
[[ 0 219  0  0  0  0  0  0  0  0]
 [ 0 287  0  0  0  0  0  0  0  0]
 [ 0 276  0  0  0  0  0  0  0  0]
 [ 0 254  0  0  0  0  0  0  0  0]
 [ 0 275  0  0  0  0  0  0  0  0]
 [ 0 221  0  0  0  0  0  0  0  0]
 [ 0 225  0  0  0  0  0  0  0  0]
 [ 0 257  0  0  0  0  0  0  0  0]
 [ 0 242  0  0  0  0  0  0  0  0]
 [ 0 244  0  0  0  0  0  0  0  0]]
```


2) Adam:

```

Epoch 19/20
125/125 [=====] - 0s 2ms/step - loss: 1.0153 - accuracy: 0.7490 - val_loss: 0.9721 - val_accuracy: 0.7720
Epoch 20/20
125/125 [=====] - 0s 1ms/step - loss: 0.9805 - accuracy: 0.7635 - val_loss: 0.9378 - val_accuracy: 0.7740
79/79 [=====] - 0s 1ms/step - loss: 1.0832 - accuracy: 0.7064

```



Test Accuracy is : 70.640

Train Accuracy is : 76.350

Train Confusion Matrix

```

[[484  0  3  11  3  14  6  0  1  0]
 [  0 548  5  3  0  4  0  2  2  0]
 [ 10 32 399 28  8  0 16  3  7  0]
 [  4 22 17 461  2  7  3 16  3  4]
 [  2  7  0  0 377  0 12  5  0 42]
 [ 23 29  7 113 22 189 11 15  9 11]
 [  7 17 14  0  5  4 465  0  1  0]
 [  3 31  1  1 10  1  0 451  1 21]
 [ 25 64 23 99 14 15 12  6 209 15]
 [ 12  9  2  6 104  0  1 109  1 239]]

```

Test Confusion Matrix

```

[[203  0  0  4  1  6  5  0  0  0]
 [  0 279  1  4  0  0  2  0  1  0]
 [ 10 20 195 28  3  0 13  7  0  0]
 [  0  8 10 213  0  2  1 15  2  3]
 [  5  7  1  0 207  0  8  6  0 41]
 [ 14  3  7 81 10 67  6 20  4  9]
 [  7  9  6  1  7  3 192  0  0  0]
 [  1 24  2  3  4  0  0 207  3 13]
 [ 10 20 16 69 13  0  2 12 90 10]
 [  4  8  1  6 73  1  0 36  2 113]]

```

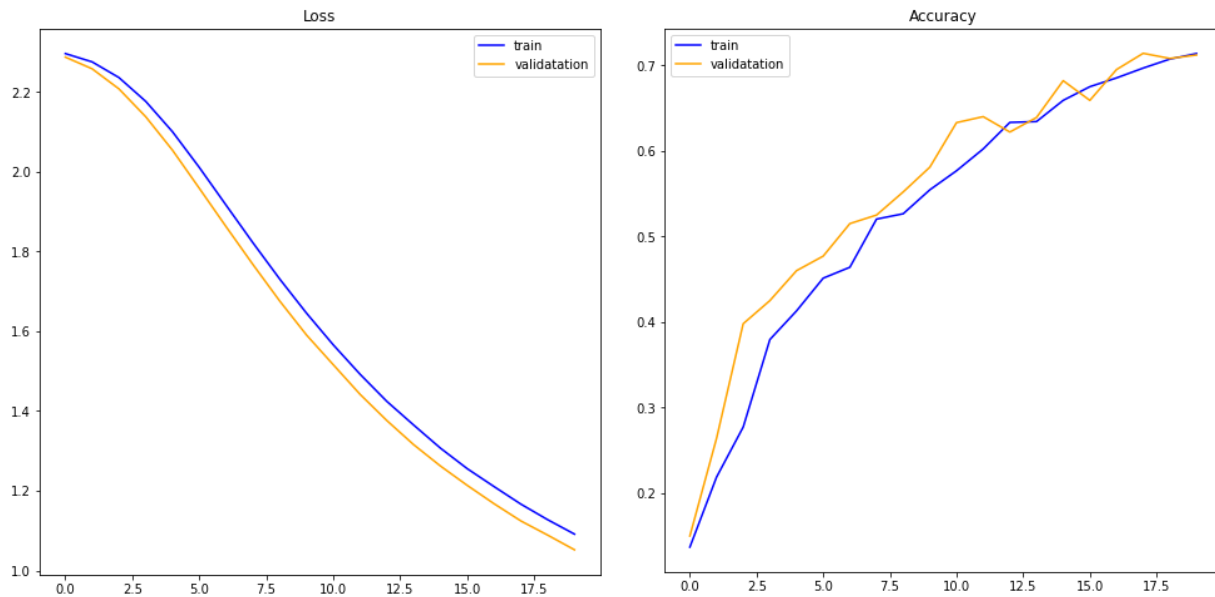
Here we tried different Optimization algorithms including SGD and ADAM. SGD could not fit on the dataset at all and as its confusion matrix shows all the output was predicted to be the class 2. This is because it could not be learned at all and the class 2 had the most elements between all.

Adam is a newer and better optimization algorithm which nowadays, developers use a lot. It is based on RMSProp activation and uses all the benefits of the optimizers. Adam is a common used optimizer and generally will result in good accuracy as it can be seen here.

4. Early Stopping:

1) True:

```
Epoch 19/20
125/125 [=====] - 0s 2ms/step - loss: 1.1275 - accuracy: 0.7072 - val_loss: 1.0888 - val_accuracy: 0.7080
Epoch 20/20
125/125 [=====] - 0s 1ms/step - loss: 1.0913 - accuracy: 0.7138 - val_loss: 1.0520 - val_accuracy: 0.7120
79/79 [=====] - 0s 1ms/step - loss: 1.1888 - accuracy: 0.6632
```



Test Accuracy is : 66.320

Train Accuracy is : 71.375

Train Confusion Matrix

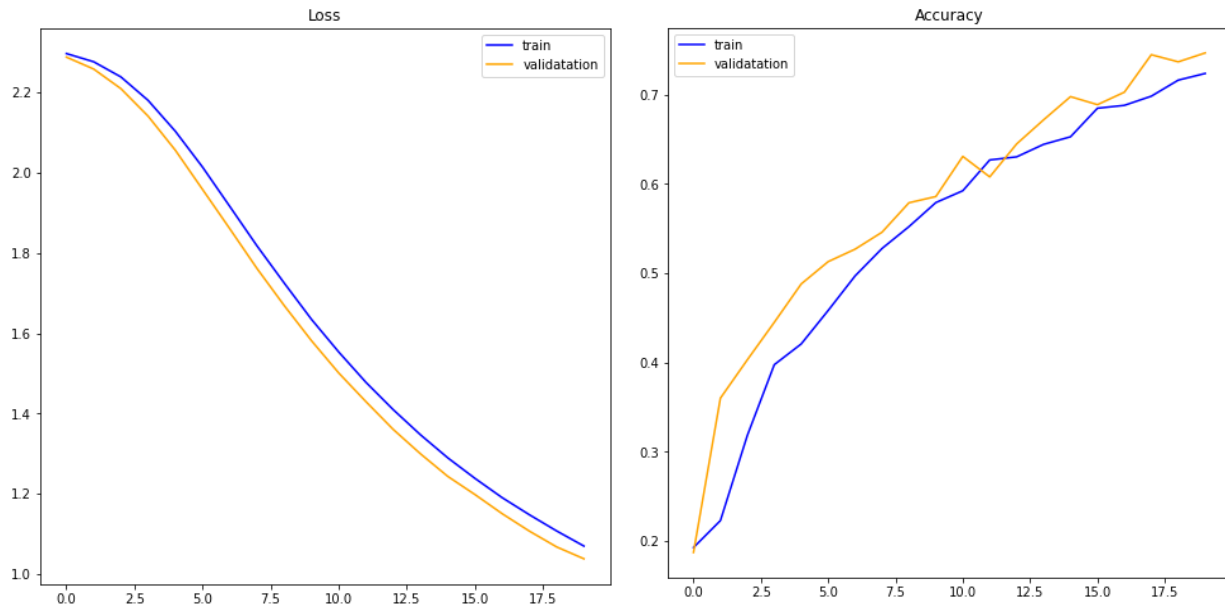
```
[[470  0  5 16  4  8 11  0  8  0]
 [ 0 551  2  5  0  0  0  2  4  0]
 [ 9 47 368 30 11  0 29  3  6  0]
 [ 4 31 15 454  1  1  3 18  8  4]
 [ 0 13  1  0 301  0 12 10  0 108]
 [18 43  5 138 21 46  7 20 117 14]
 [ 6 25 19  0  3  1 454  0  5  0]
 [ 4 42  0  0  4  0  1 458  1 10]
 [22 88 15 94 10  1 11 14 214 13]
 [12 13  1  5 62  0  2 162  1 225]]
```

Test Confusion Matrix

```
[[193  0  0  8  2  3 10  0  3  0]
 [ 0 282  2  2  0  0  1  0  0  0]
 [ 8 39 173 28  2  0 18  7  1  0]
 [ 0 10 10 213  0  0  1 16  3  1]
 [ 2  9  1  0 161  0 13 12  0 77]
 [ 8  6  6 98  5 21  6 23 38 10]
 [ 6 14  7  1  5  2 189  0  1  0]
 [ 1 31  2  1  3  0  0 209  2  8]
 [ 8 33 10 62  9  0  2 17 90 11]
 [ 2 12  1  4 42  0  2 52  2 127]]
```

2) False:

Epoch 19/20
 125/125 [=====] - 0s 2ms/step - loss: 1.1076 - accuracy: 0.7165 - val_loss: 1.0678 - val_accuracy: 0.7370
 Epoch 20/20
 125/125 [=====] - 0s 1ms/step - loss: 1.0700 - accuracy: 0.7240 - val_loss: 1.0383 - val_accuracy: 0.7470
 79/79 [=====] - 0s 1ms/step - loss: 1.1730 - accuracy: 0.6792



Test Accuracy is : 67.920

Train Accuracy is : 72.400

Train Confusion Matrix

```
[[471  0  3  4  4 11 11  1 17  0]
 [  0 550  3  1  0  0  0  1  8  1]
 [ 11 45 342 31 19  0 39  3 13  0]
 [  6 29 22 404  4  3  2 17 45  7]
 [  1  9  0  0 354  0 11  3  0 67]
 [ 21 37 12 72 31 55 14 12 154 21]
 [  7 23  3  0  6  1 469  0  4  0]
 [  2 34  0  0 11  0  0 430  2 41]
 [ 19 68 14 36 16  1 13  6 290 19]
 [ 12 11  1  1 81  0  1  67  6 303]]
```

Test Confusion Matrix

```
[[192  0  0  3  2  5 11  0  6  0]
 [  0 281  1  1  0  0  2  0  2  0]
 [  9 33 164 21  4  0 29  7  8  1]
 [  0  8 18 187  0  0  1 17 18  5]
 [  0  8  1  0 197  0 10  3  0 56]
 [ 13  3 11 54 13 14  7 20 76 10]
 [  6 12  2  1 10  1 191  0  2  0]
 [  1 26  2  0  4  0  0 196  3 25]
 [ 10 23 13 23 11  0  3 11 129 19]
 [  4  8  1  2 58  0  0 20  4 147]]
```

Here we tried Early Stopping.

In many cases our model will stop training after it reach a converge or a local or global minimum. After that it should be stopped since learning on the same data with little update makes the model overfitted to that dataset. So, after a couple of epochs when there are small changes in loss our model should be stopped to prevent overfitting.

In this question, unfortunately, we could not see the early stopping result since it could not reach a convergence after 20 epoch and loss value still had big changes. To see the results, we should have increased the epochs, but we know the result, after few epochs in which loss had small changes, our model will stop the training to prevent overfitting.

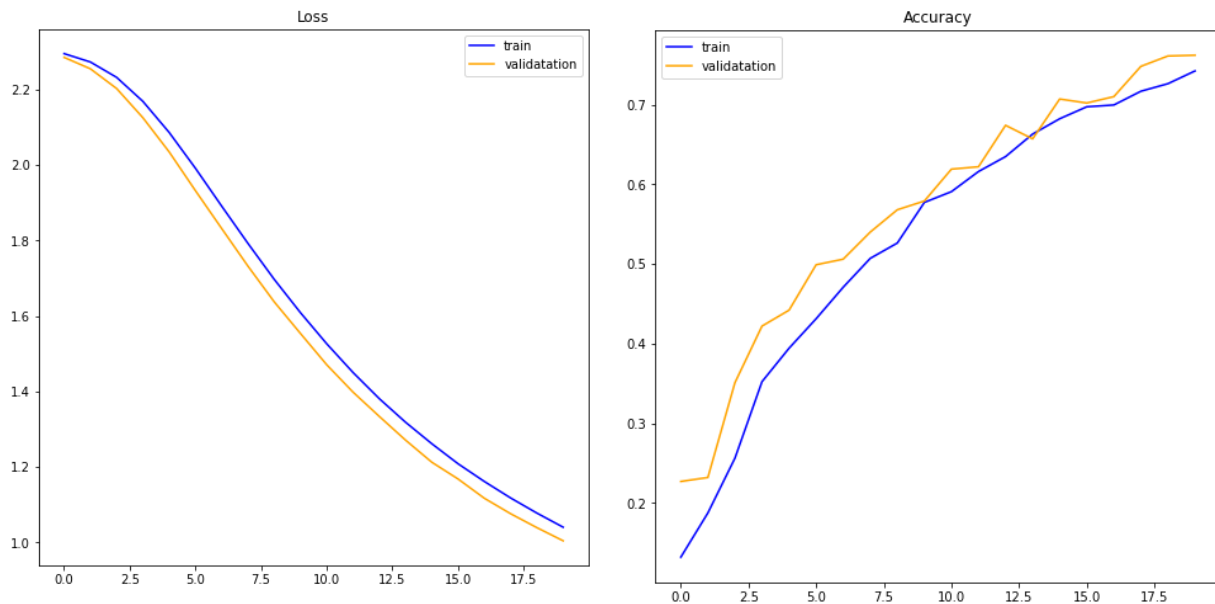
5. Epochs:

1) 20 Epoch:

```

...
Epoch 19/20
125/125 [=====] - 0s 2ms/step - loss: 1.0779 - accuracy: 0.7262 - val_loss: 1.0393 - val_accuracy: 0.7
610
Epoch 20/20
125/125 [=====] - 0s 2ms/step - loss: 1.0405 - accuracy: 0.7423 - val_loss: 1.0044 - val_accuracy: 0.7
620
79/79 [=====] - 0s 1ms/step - loss: 1.1435 - accuracy: 0.6952

```



Test Accuracy is : 69.520

Train Accuracy is : 74.225

Train Confusion Matrix

```

[[489  0  1  3  2 14 12  0  1  0]
 [  0 549  4  4  0  4  0  1  1  1]
 [ 13 36 371 26  8  1 36  3  8  1]
 [  6 25 15 448  2  9  5 19  4  6]
 [  2  7  0  0 311  0 15  7  0 103]
 [ 26 24  4 106 23 192 20 21  0 13]
 [  7 17  2  0  3  5 479  0  0  0]
 [  4 27  0  0  5  0  0 447  3 34]
 [ 24 65 25 96  8 16 20 11 198 19]
 [ 12  8  1  4 50  0  2 121  1 284]]

```

Test Confusion Matrix

```

[[203  0  0  0  0  6 10  0  0  0]
 [  0 280  0  4  0  0  3  0  0  0]
 [ 11 25 169 21  2  1 30  9  7  1]
 [  1  8  6 205  0  7  3 18  2  4]
 [  6  8  0  0 167  0 10  8  0 76]
 [ 15  3  6 73  7 70 11 26  1  9]
 [  7  7  0  1  6  2 202  0  0  0]
 [  1 21  1  1  3  0  0 203  4 23]
 [ 10 22 15 58  9  9  3 16 83 17]
 [  4  6  1  3 37  2  0 33  2 156]]

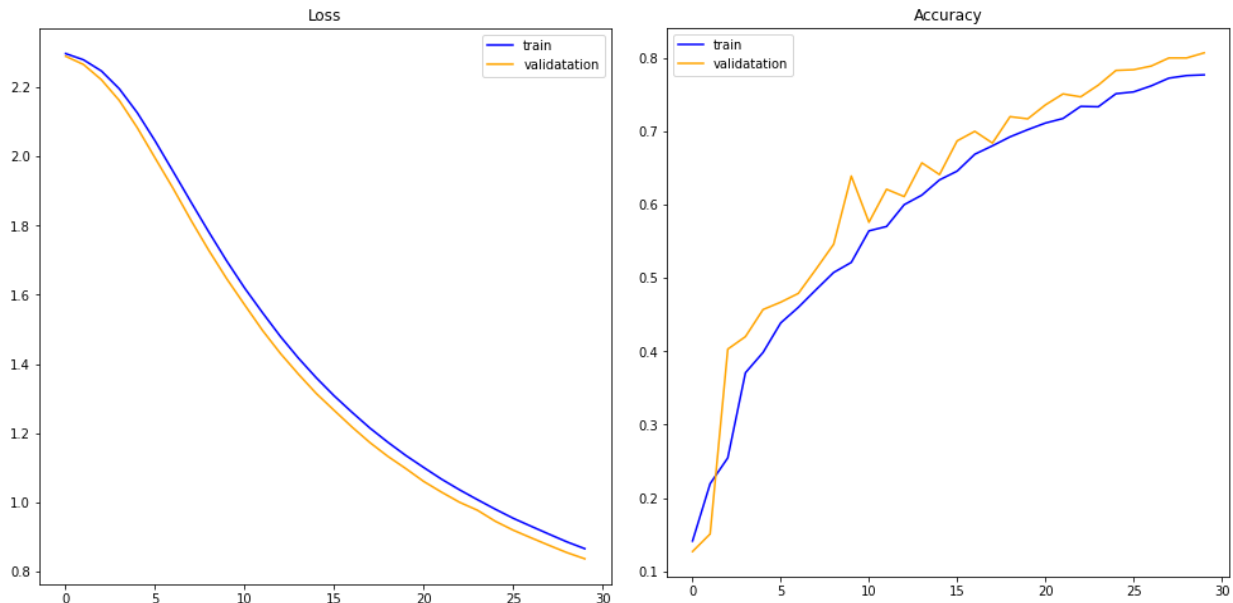
```

2) 30 Epoch:

```

Epoch 29/30
125/125 [=====] - 0s 2ms/step - loss: 0.8859 - accuracy: 0.7760 - val_loss: 0.8549 - val_accuracy: 0.8000
Epoch 30/30
125/125 [=====] - 0s 2ms/step - loss: 0.8662 - accuracy: 0.7770 - val_loss: 0.8370 - val_accuracy: 0.8070
79/79 [=====] - 0s 1ms/step - loss: 0.9802 - accuracy: 0.7364

```



Test Accuracy is : 73.640

Train Accuracy is : 77.700

Train Confusion Matrix

```

[[488  0  2  3  2 16  8  0  3  0]
 [ 0 544  7  3  0  2  0  2  6  0]
 [10 19 406 24 10  0 23  3  8  0]
 [ 4 14 21 446  1 13  3 16 14  7]
 [ 2  4  1  0 357  0 11  4  1 65]
 [24 14  9 89 21 217 12 14 19 10]
 [ 7 10  6  0  4  3 481  0  2  0]
 [ 4 19  2  1  8  1  0 442  5 38]
 [20 34 27 54 16 10 11  3 288 19]
 [12  7  2  5 84  0  1 66  3 303]]

```

Test Confusion Matrix

```

[[208  0  0  0  1  6  4  0  0  0]
 [ 0 279  1  1  0  0  2  0  4  0]
 [ 8 11 202 17  2  1 23  7  4  1]
 [ 1  4 13 203  0  7  1 12  8  5]
 [ 1  5  1  0 198  0  9  2  0 59]
 [15  0  8 65 10 79  7 22  7  8]
 [ 7  3  4  0  5  3 203  0  0  0]
 [ 2 15  7  1  4  0  0 197  3 28]
 [10 12 18 40  9  2  2  9 125 15]
 [ 4  5  1  3 53  2  0 24  5 147]]

```

Here we are comparing the result of different Epochs.

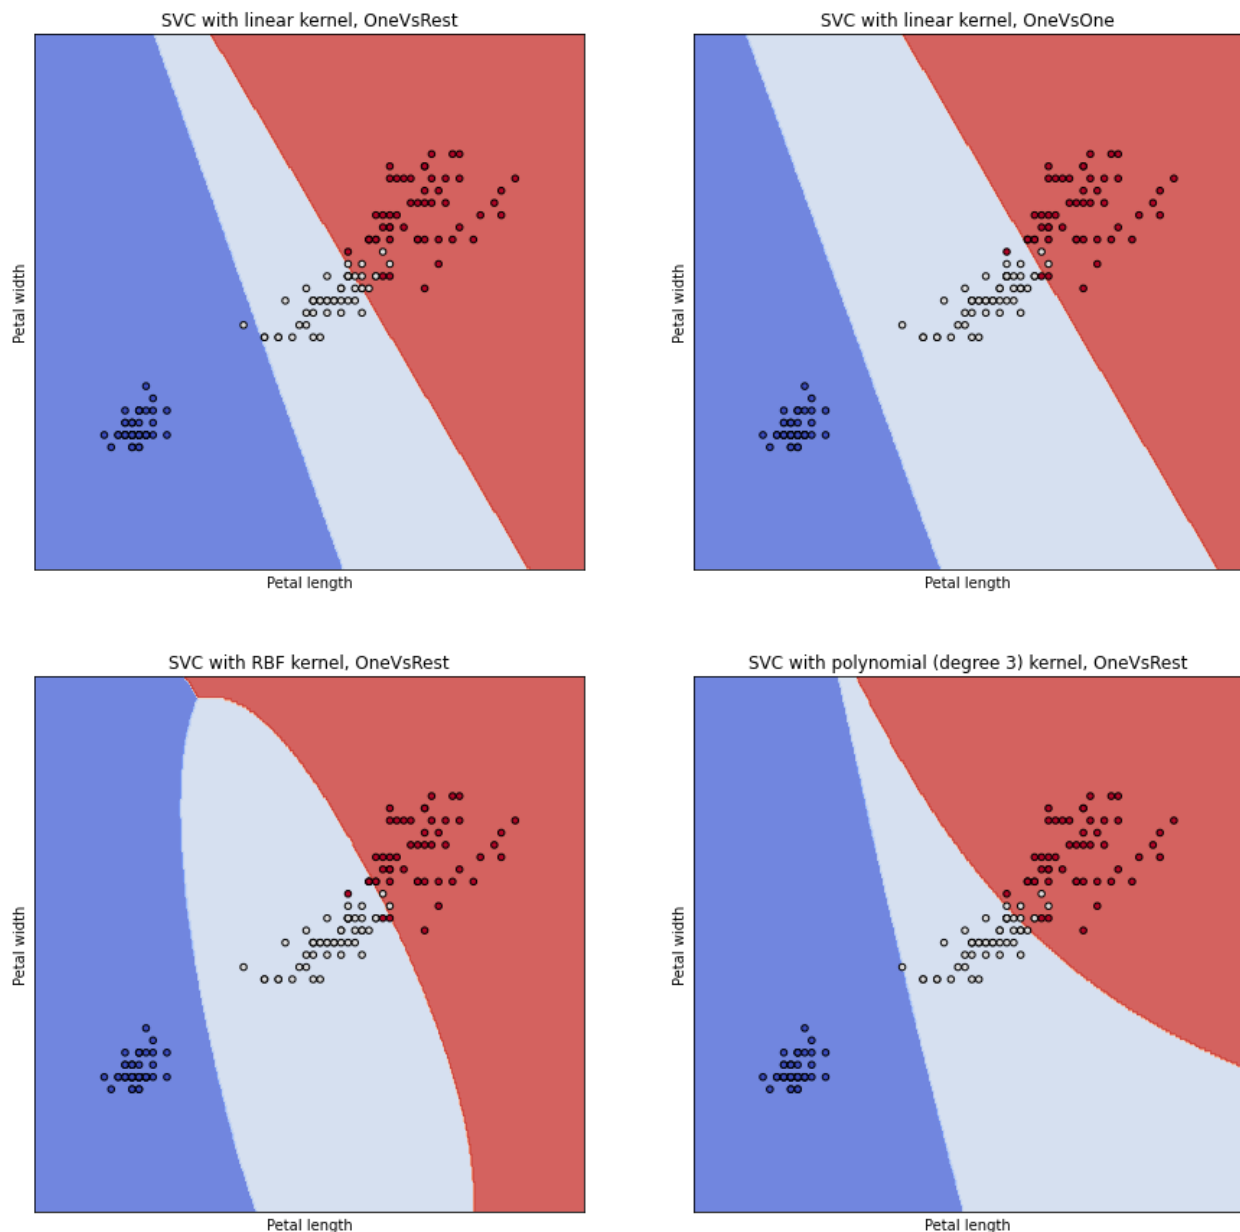
As it can be seen, for 30 epochs we had a better accuracy. This is because after 20 epochs our model did not reach a converge and it was still learning, so more epochs were needed to learn and reach a minimum.

After 30 epochs, we can see our model still did not reached a minimum but it is going closer to it. We still need more epochs for the model to be trained perfectly but as we can see 30 epochs reached a better accuracy than 20 because of the fact mentioned above.

Question 9: Coding:

In this Question we are going to use Petal Length and Petal Width features of Iris dataset to Classify the dataset using different SVM kernels using different methods including (OneVsRest and OneVsOne). Then we are going to plot the areas divided using SVM classifier along with the accuracy and confusion matrix of each part:

To compare the kernels area we have the plot below:

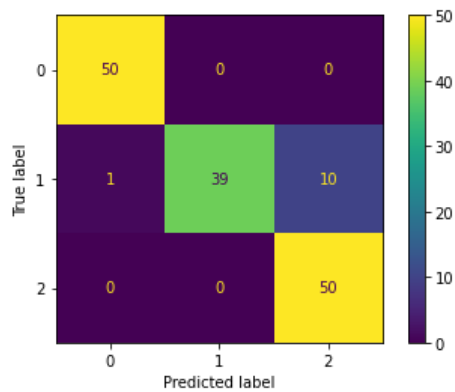


For Each one we have:

1)

SVC with linear kernel, OneVsRest

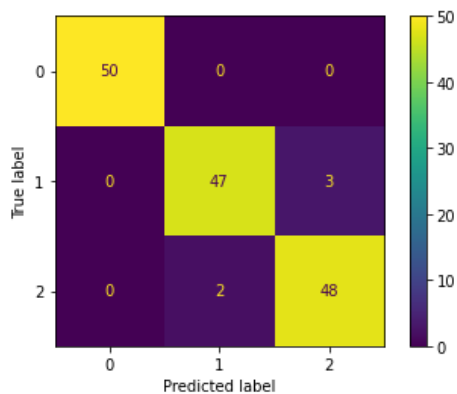
	precision	recall	f1-score	support
0	0.98	1.00	0.99	50
1	1.00	0.78	0.88	50
2	0.83	1.00	0.91	50
accuracy			0.93	150
macro avg	0.94	0.93	0.93	150
weighted avg	0.94	0.93	0.93	150



2)

SVC with linear kernel, OneVsOne

	precision	recall	f1-score	support
0	1.00	1.00	1.00	50
1	0.96	0.94	0.95	50
2	0.94	0.96	0.95	50
accuracy			0.97	150
macro avg	0.97	0.97	0.97	150
weighted avg	0.97	0.97	0.97	150



3)

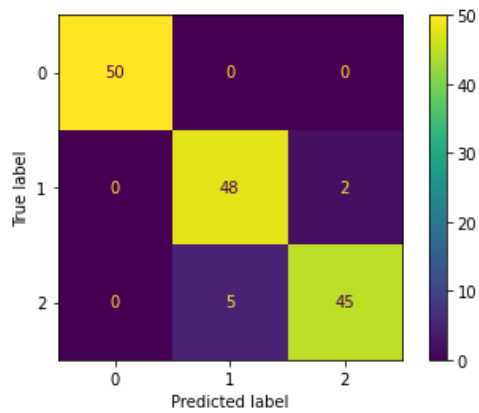
```

SVC with RBF kernel, OneVsRest
      precision    recall  f1-score   support

     0         1.00      1.00      1.00        50
     1         0.91      0.96      0.93        50
     2         0.96      0.90      0.93        50

 accuracy         0.95         0.95         0.95        150
 macro avg         0.95         0.95         0.95        150
 weighted avg         0.95         0.95         0.95        150

```



4)

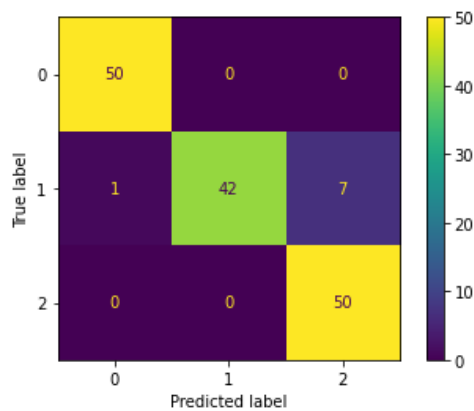
```

SVC with polynomial (degree 3) kernel, OneVsRest
      precision    recall  f1-score   support

     0         0.98      1.00      0.99        50
     1         1.00      0.84      0.91        50
     2         0.88      1.00      0.93        50

 accuracy         0.95         0.95         0.95        150
 macro avg         0.95         0.95         0.95        150
 weighted avg         0.95         0.95         0.95        150

```



#####

We can see OneVsOne Linear kernel achieved best accuracy. Even in the first plot it is obvious that Linear with One vs One had the lowest error in classification. All the confusion matrix plots and accuracies are shown. Linear Kernel Using one vs one method reached an accuracy of 97% which was the highest of all.

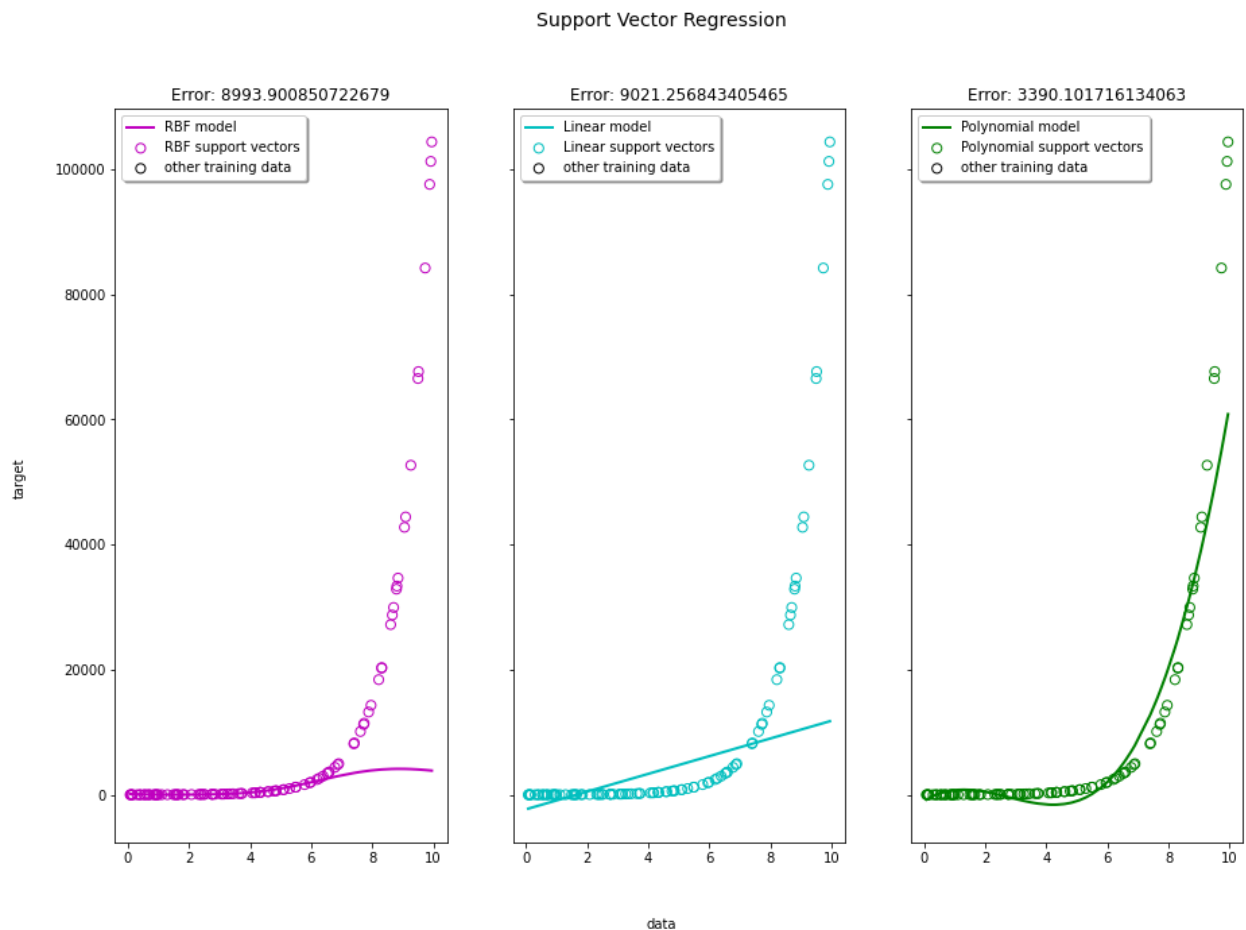
Question 10: Coding:

1) S

Here first we are going to produce 100 point between 0 and 10 with respect to:

$$y = 5e^x + 3; x \in [0, 10]$$

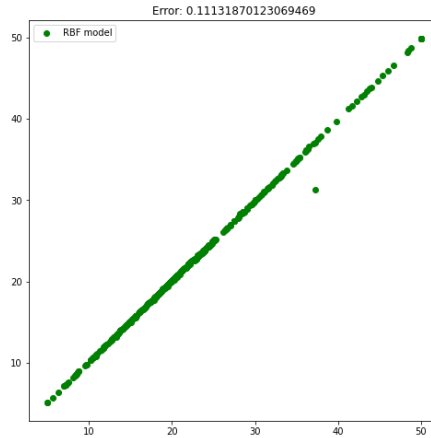
And then run 3 SVM regression with polynomial third degree, rbf, and exponential kernels. Our result is like below:



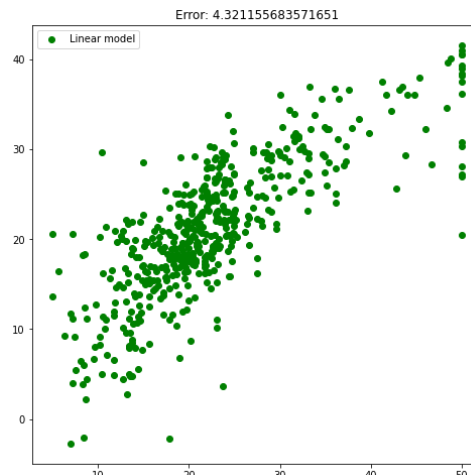
Errors are come as title for each plot. Polynomial third degree had the best regressor with lowest error among all.

- 2) Here we do the same thing as in part 1 but on the Boston housing dataset. We are going to predict the price of each house using a svm regression with rbf, polynomial third degree, and exponential kernels. Our results are:

For RBF:



For Linear:



For Polynomial with degree 3: It took a lot of time solving a polynomial problem. So, results for this kernel could not be shown.

SVM for a dataset consisting many samples is not a good idea. Since SVM should solve a quadratic optimization problem it took a lot of time for rbf or polynomial to be solved. So, with Boston housing dataset, it took a lot of time fitting a polynomial third degree on the dataset. But as it can be seen a RBF kernel SVR fitted well.

All Errors are Mean Absolute Error which are on the title for each plot.